

Alvin Christian Nataputra

Jakarta, Indonesia • alvincnataputra@gmail.com • github.com/apricitea • apricitea.github.io • linkedin.com/in/alvincnataputra

Research interests

Reliable machine learning systems: validity and leakage control in time-series ML, distribution shift, and how autonomous agents stay grounded over long horizons. Applied work on Indonesian-language text and Indonesian equity markets, where data is thin and regimes change.

Education

IPB University (Bogor Agricultural University)

Bogor, Indonesia

Bachelor of Statistics and Data Science • August 2019 – August 2023 • GPA 3.83/4.00

Thesis: suicide-ideation detection in Indonesian user-generated text. Built a FastText + LSTM classifier and benchmarked it against a fine-tuned IndoBERT transformer on an identical 80/20 split of 10,128 labelled posts: positive-class F1 rose from 0.79 to 0.90 and accuracy from 0.91 to 0.96. Class-imbalance treatment (class weighting, ADASYN) *reduced* F1 in the LSTM arm — a negative result reported rather than dropped. The model is deployed as a public demo; the work is research/educational and is not a clinical tool.

Research

Reproducible evaluation of an IDX equity signal pipeline

2025–2026

Independent research; ensemble ML for the Indonesia Stock Exchange

- Audited an existing evaluation of a LightGBM/XGBoost signal pipeline and found four defects that invalidated its reported performance: the final model was trained on rows spanning the test window, walk-forward early stopping used the fold's own test set, the feature scaler was fitted on the whole series before splitting, and the embargo was five calendar days against a ten-bar label horizon with no purging of overlapping labels.
- Rebuilt the evaluation with purged and embargoed folds, fold-local scaling, average-uniqueness sample weights for overlapping labels, and buy-and-hold and momentum baselines under identical IDX transaction costs.
- Reported result is negative and is the point of the work: the model returns +0.12% / +7.62% / +4.73% across three configurations, against +14.17% for buy-and-hold and +20.71% for the index; walk-forward accuracy of 0.34–0.36 sits below the 0.53 majority-class baseline. The meta-labeler claimed to improve risk-adjusted return instead reduced exposure and most of the return.
- Committed a full run manifest — per-file data SHA-256, exact split dates, seeds, configuration, environment lockfile, trade ledger and predictions — so the result is re-derivable rather than asserted.
github.com/apricitea/project-aurum/tree/main/research

Bahasa Indonesia corpus: collection pipeline and rights-cleared release

2026

Public dataset construction with a provenance gate

- Built and operate a weekly pipeline collecting Indonesian text (Wikipedia extracts and news sources), with an explicit release rule: a record is published only if its source maps to a licence that can be asserted *and* it carries source URL, revision ID, title and retrieval timestamp.
- Released 552 attributable records under CC BY-SA 4.0 with per-record attribution; excluded roughly 2,200 news-derived and 30,000 unattributed records rather than publish them without a licence.
huggingface.co/datasets/apricitea/bahasa-indonesia-corpus

Market-regime dating for IDX equities

2024

Implemented a censored local-extrema algorithm in the style of Pagan & Sossounov: candidate turning points are filtered by strict peak/trough alternation and phases must clear minimum duration and amplitude thresholds, so no arbitrary fixed lookback window is assumed. Validated by a test suite and shown to isolate the 2020 COVID drawdown as a distinct bear phase on BBRI without supervision.

Causal effect estimation in telco churn

2023

Extended an XGBoost churn model (SHAP-explainable) with an inverse-probability-of-treatment-weighting estimator and bootstrapped confidence intervals, to separate correlation from causation for product-adoption decisions. Call Center adoption causally reduced churn and increased CLTV; video and MyApp adoption increased monthly purchase but their CLTV effects were indistinguishable from zero — which changes the recommendation from “push adoption” to “fix the product first”.

Telecommunications CVM analytics and ML lab

2026

Idempotent DuckDB warehouse with a data-quality gate that aborts before touching downstream tables, two independently calibrated risk models (churn AUC 0.823; late-payment AUC 0.628), and a test that enforces the forward-looking construction of the late-payment label.

Experience

XLSmart — Data Scientist, Customer Value & Experience

Aug 2024 – present

Jakarta, Indonesia • formerly XL Axiata • third-largest Indonesian telco, 80M+ subscribers

- Data science and analytics partner for XL Prioritas (postpaid), Customer Experience and Go-to-Market teams; own end-to-end analytics from opportunity identification through segmentation, targeting and impact measurement.
- Built churn-driver and high-risk-segment analysis that shaped retention initiatives; the resulting programme was reported to improve customer retention by 25%.
- Led end-to-end analytics for targeted campaigns, enabling commercial teams to optimise targeting and measure incremental revenue.
- Built an agentic-AI analytics layer used by senior leadership to monitor key metrics, identify performance drivers and receive recommendations.

Stack: Python, SQL, BigQuery, Databricks, Snowflake, Tencent DLC/WeData.

Asta Protek Jiarsi — Data Scientist, IT consulting

Sep 2023 – Jun 2024

Jakarta, Indonesia • consulting and compliance services

- Delivered ML models, segmentations and dashboards for anti-money-laundering, fraud detection and enterprise analytics, working with client teams to production.
- Standardised ETL workflows and Docker-based ML deployment pipelines, removing a recurring class of production incidents and improving reproducibility.
- Applied SHAP to explain model predictions, supporting business interpretation and regulatory compliance for client ML solutions.

Stack: Python, Poetry/uv, PostgreSQL, SQL Server, Docker.

Indosat Ooredoo Hutchison — Data Science & Analytics Intern

May – Aug 2023

Jakarta, Indonesia • second-largest Indonesian telco, 100M+ subscribers

Automated recurring business analyses from spreadsheets into dashboards and scheduled reports, cutting manual reporting workload by 20%; ran deep-dive statistical testing on product and campaign performance, identifying Pareto contributors.

Bank Rakyat Indonesia (BRI) — Analytics Intern

Oct 2022 – Apr 2023

Jakarta, Indonesia • state-owned bank

Built a non-performing-loan risk model reaching 60%+ recall at 70% precision; produced geospatial analyses supporting regional coverage optimisation (10%+ operational efficiency improvement).

Technical

Languages Python, SQL, R, shell

Machine learning gradient-boosted trees (LightGBM, XGBoost), calibrated classification, clustering, feature engineering, model evaluation, explainability (SHAP), transformers (BERT/IndoBERT), LSTM, FastText

Statistics & experimentation A/B testing, control-group design, hypothesis testing, causal inference (IPTW, propensity scores), time-series validation (walk-forward CV, purging, embargo)

Engineering & reproducibility Docker, uv, pytest, GitHub Actions CI, Git, PostgreSQL, DuckDB, Qdrant

Data platforms BigQuery, Databricks, Snowflake, Tencent WeData/DLC

Natural languages Indonesian (native), English (Cambridge C1 Advanced in preparation; examination targeted 2027)

Service

Volunteer computing on a self-hosted homelab: 247,311 BOINC credits across Einstein@Home (gravitational-wave and pulsar analysis), MilkyWay@Home and World Community Grid. Also operate a Tor Snowflake proxy supporting censorship circumvention.